

# **Analysis of the Hugging Face Transformers Library: Purpose and Component Classes**

## **1. Introduction to Modern Machine Learning and the Role of Specialized Libraries**

The field of machine learning has witnessed remarkable advancements, leading to the development of increasingly complex models capable of tackling intricate tasks across diverse domains. This sophistication often necessitates the creation of specialized libraries that provide developers and researchers with pre-built tools and functionalities. These libraries abstract away the underlying complexities, allowing practitioners to focus on applying cutting-edge techniques to specific problems without needing to implement every component from scratch. A significant trend in modern machine learning is the utilization of pre-trained models. These models, having been trained on vast amounts of data, learn general-purpose features that can be effectively transferred to new, often smaller, datasets and specific tasks. This approach significantly reduces the computational resources and time required for training while often yielding superior performance, particularly in areas such as Natural Language Processing (NLP), Computer Vision, and Audio processing, where obtaining large, labeled datasets can be challenging. This report will delve into a prominent Python library that embodies this paradigm, offering a wide array of pre-trained models and the necessary tools to leverage them effectively across these domains.

## **2. Identifying the Core Library: Hugging Face Transformers**

An examination of the provided class names reveals several key patterns. Prefixes such as 'AST', 'Albert', 'Align', 'AltCLIP', and 'Aria' likely denote specific model architectures or families of models. The consistent presence of suffixes like 'Config', 'Model', 'Tokenizer', and 'Processor' strongly suggests a comprehensive library designed to handle various transformer-based architectures and their associated utilities. These suffixes typically indicate the role of each class within the library's structure, with 'Config' classes managing model configurations, 'Model' classes representing the neural network architectures themselves, 'Tokenizer' classes handling the conversion of text into numerical inputs, and 'Processor' classes dealing with data preparation for different modalities.

Further investigation through the provided research material confirms the identity of this library. Several snippets explicitly mention "Hugging Face Transformers" (or its earlier names like "pytorch-transformers") and directly associate it with the types of

classes listed in the query.<sup>1</sup> For instance, one snippet directly answers the query by stating that classes like `ASTConfig`, `AlbertModel`, and `AutoTokenizer` are part of a Python library featuring "Auto Classes".<sup>1</sup> Another snippet explicitly lists these "AutoModels" as belonging to the transformers library.<sup>2</sup> The official documentation of 🤗 Transformers, as seen in multiple snippets, consistently refers to classes such as `AutoConfig`, `AutoModel`, and `AutoTokenizer` as integral components.<sup>3</sup> Mentions of "transformers" in the context of GitHub issues also indicate an active and evolving library within the machine learning community.<sup>10</sup> The recurring and direct association of these class types with the name "Hugging Face Transformers" across various reliable sources leaves little doubt about the library in question. Therefore, it can be definitively stated that the Python library containing the listed class names is the Hugging Face Transformers library.

### **3. Purpose and Key Features of the Hugging Face Transformers Library**

The official documentation provides a clear understanding of the Hugging Face Transformers library's purpose. It serves as a comprehensive platform for state-of-the-art machine learning, primarily focused on pre-trained models for PyTorch, TensorFlow, and JAX.<sup>3</sup> The library offers an extensive collection of thousands of pre-trained models that can be readily applied to various tasks across different modalities, including text, vision, and audio.<sup>3</sup> These tasks range from fundamental ones like text and image classification to more complex applications such as question answering, summarization, translation, text and image generation, object detection, speech recognition, and even multimodal tasks that combine different types of data.<sup>3</sup>

A key feature of the library is its support for interoperability between popular deep learning frameworks. Users can seamlessly switch between PyTorch, TensorFlow, and JAX at different stages of a model's lifecycle, allowing for flexibility in choosing the most suitable framework for training, evaluation, and deployment.<sup>3</sup> The library provides a unified and user-friendly API for accessing and utilizing these diverse models. This ease of use is further emphasized by the availability of high-level abstractions like the pipeline function and 'Auto' classes, which simplify common tasks and model loading.<sup>3</sup> Beyond simply providing pre-trained models, the library also offers robust tools for fine-tuning these models on custom datasets, enabling users to adapt them to their specific needs.<sup>2</sup> Utilities for essential data processing steps, such as tokenization (through `AutoTokenizer`) and feature extraction/image processing (through `AutoFeatureExtractor` and `AutoImageProcessor`), are also integral parts of the library.<sup>1</sup> Furthermore, the Hugging Face Transformers library is tightly integrated with the Hugging Face Model Hub, a central repository where users can discover, share, and collaborate on models, datasets, and even machine learning

applications.<sup>3</sup> This ecosystem fosters collaboration and makes advanced machine learning resources readily available to a wide audience. By providing access to pre-trained models, the library significantly reduces the computational costs, carbon footprint, and time associated with training complex models from scratch.<sup>3</sup> The overall aim of the Hugging Face Transformers library is to make state-of-the-art machine learning accessible to a broader community by offering a powerful yet user-friendly platform for utilizing and adapting pre-trained transformer models across a multitude of tasks and modalities. This focus on accessibility and collaboration has contributed to its widespread adoption in both research and practical applications.

## **4. Deconstructing the Class Names: A Categorical Overview**

To gain a deeper understanding of the Hugging Face Transformers library, the provided class names can be organized into functional categories, revealing the library's underlying architecture and the roles of its various components.

### **4.1. Configuration Classes (Suffix: Config)**

Configuration classes, such as `ASTConfig`, `AlbertConfig`, and `AlignConfig`, serve as blueprints for defining the architecture and hyperparameters of the models within the library.<sup>1</sup> These classes store essential settings that govern the model's structure and behavior, including the number of layers, the size of hidden dimensions, the number of attention heads, and various other architectural parameters. The `PretrainedConfig` class acts as a base for these specific model configurations, providing common attributes and functionalities.<sup>13</sup> The library also offers the `AutoConfig` class, which simplifies the process of loading configurations by automatically retrieving the appropriate configuration class based on a given pre-trained model name or path.<sup>1</sup> This automated loading mechanism, as illustrated in the library's internal structure<sup>14</sup>, streamlines the workflow for users as they do not need to manually specify the configuration class for each model they wish to use. Furthermore, the library allows for customization of these configurations by subclassing `PretrainedConfig`, enabling users to tailor model architectures to their specific requirements.<sup>15</sup> These configuration classes are crucial for instantiating models with the correct architecture and for ensuring compatibility with pre-trained weights. For instance, an `AlbertConfig` object would encapsulate all the specific architectural details of an Albert model, ensuring that when an Albert model is created, it adheres to the intended design.

### **4.2. Model Classes (Suffix: Model, For...)**

The core of the Hugging Face Transformers library lies in its model classes, which

implement the actual neural network architectures. These classes can be broadly divided into base model classes and task-specific model classes. Base model classes, such as `ASTModel`, `AlbertModel`, and `AlignModel`, provide the foundational transformer architecture.<sup>1</sup> Task-specific model classes, on the other hand, extend these base models by adding specialized layers or "heads" designed for particular machine learning tasks. Examples from the list include `AlbertForMaskedLM` (for masked language modeling), `AlbertForSequenceClassification` (for classifying sequences), and `AutoModelForAudioClassification` (for classifying audio).<sup>1</sup> The `AutoModel` class acts as a versatile entry point, automatically loading the appropriate base model class based on a pre-trained model identifier.<sup>1</sup> Similarly, there are task-specific '`AutoModelFor...`' classes that automatically select the model with the correct task-specific head.<sup>2</sup> The distinction between `AutoModel` and task-specific variants is that the latter includes additional layers tailored for a specific objective, such as a classification layer for `AutoModelForSequenceClassification`.<sup>18</sup> The library's design, with base models and task-specific extensions, offers a modular approach, allowing users to easily leverage pre-trained architectures for a wide range of downstream tasks. The `PreTrainedModel` class serves as a base for all model classes, providing common methods for loading, saving, and manipulating models.<sup>16</sup> For instance, while `AlbertModel` provides the core Albert transformer architecture, `AlbertForQuestionAnswering` builds upon this by adding a question answering head to predict the start and end tokens of an answer within a given context.

#### **4.3. Tokenizer Classes (Suffix: `Tokenizer`, `TokenizerFast`)**

Tokenizer classes are essential for processing text data in the Hugging Face Transformers library. They are responsible for converting raw text into numerical input IDs that can be fed into the transformer models. This process involves several steps, including splitting the text into smaller units called tokens, converting these tokens into numerical IDs based on a vocabulary, and adding special tokens that the model may rely on.<sup>22</sup> The library provides various tokenizer classes, such as `AlbertTokenizer` and `AlbertTokenizerFast`, each tailored to work with specific model architectures.<sup>1</sup> A key distinction exists between tokenizers ending in `Tokenizer` and those ending in `TokenizerFast`. The 'Fast' variants, like `AlbertTokenizerFast`, are typically implemented in Rust, offering significant speed advantages, especially when processing large amounts of text or performing batched tokenization.<sup>22</sup> They also often provide additional functionalities, such as advanced alignment methods between the original text and the tokenized output.<sup>23</sup> The `AutoTokenizer` class acts as a convenient way to automatically load the appropriate tokenizer for a given pre-trained model.<sup>1</sup> This simplifies the user experience as the library handles the selection of the correct

tokenizer based on the model being used. Different models may employ different tokenization algorithms, such as Byte-Pair Encoding (BPE), WordPiece, or SentencePiece, and the corresponding tokenizer classes implement these specific algorithms.<sup>24</sup> For example, when using `AlbertTokenizerFast` with the sentence "Hello, world!", it would tokenize the sentence into subword units, convert these units into numerical IDs based on the Albert model's vocabulary, and add special tokens like at the beginning and at the end, as required by the model.

#### **4.4. Processor Classes (Suffix: `Processor`)**

Processor classes, such as `AlignProcessor`, `AltCLIPProcessor`, and `AriaProcessor`, play a crucial role in preparing data for multimodal models in the Hugging Face Transformers library.<sup>2</sup> Multimodal models are designed to handle different types of input data, such as text and images, and processor classes often encapsulate the necessary preprocessing steps for each modality. They can integrate the functionalities of both tokenizers (for text) and feature extractors or image processors (for other modalities) into a single class. While the provided snippets do not offer detailed explanations of each specific processor class, their naming convention and association with models like `AlignModel` (likely for text-image alignment) and `AltCLIPModel` (for contrastive language-image pre-training) strongly suggest their function in handling the specific preprocessing requirements for these multimodal tasks. For instance, an `AlignProcessor` might take both text and image inputs, use an internal tokenizer to process the text, and perform image transformations like resizing and normalization on the image, ultimately producing a unified set of inputs suitable for the `AlignModel`. These processor classes are essential for streamlining the data preparation pipeline for complex multimodal applications, ensuring that the input data is correctly formatted and preprocessed for the corresponding model.

#### **4.5. Feature Extractor and Image Processor Classes (Suffix: `FeatureExtractor`, `ImageProcessor`)**

For models that handle non-textual data, the Hugging Face Transformers library provides specialized classes like feature extractors and image processors. Feature extractor classes, such as `ASTFeatureExtractor`, are designed to extract relevant features from data like audio. In the case of `ASTFeatureExtractor`, it likely processes raw audio data to create a spectrogram representation, which is a visual representation of the audio's frequency spectrum over time, suitable for input to the `ASTModel` (Audio Spectrogram Transformer). Similarly, image processor classes, such as `AriaImageProcessor`, are responsible for preprocessing images before they are fed into image-based models. This preprocessing can include steps like resizing the

image to a specific resolution, normalizing pixel values, and converting the image into a tensor format that the model can understand. The library also provides `AutoFeatureExtractor` and `AutoImageProcessor` classes<sup>1</sup>, which, similar to other 'Auto' classes, automatically load the appropriate feature extractor or image processor based on the pre-trained model being used. These specialized classes extend the library's capabilities to handle a broader range of data modalities beyond just text, enabling the use of transformer models for tasks involving audio and images.

#### **4.6. Utility and Helper Classes**

The Hugging Face Transformers library includes a variety of utility and helper classes that provide supporting functionalities for various aspects of the machine learning workflow. Optimization algorithms like `Adafactor` and `AdamWeightDecay` are provided for training models by efficiently updating their weights.<sup>2</sup> `AdaptiveEmbedding` represents a type of embedding layer that can dynamically adjust its size based on token frequencies, potentially improving memory usage and performance for models with large vocabularies. The `AddedToken` class allows users to define and add special tokens to a tokenizer's vocabulary, which can be useful for handling custom or domain-specific tokens not present in the default vocabulary.<sup>23</sup> The `Agent` class might relate to more advanced applications like reinforcement learning or dialogue systems, suggesting the library's potential scope beyond standard supervised learning tasks. `AttentionInterface` likely serves as an abstract base class or interface for defining different attention mechanisms used within the transformer models. `AsyncTextIteratorStreamer` is a utility for efficiently handling and streaming large text datasets asynchronously, which can be particularly useful for tasks like text generation where long sequences are produced. These utility and helper classes demonstrate the library's comprehensive nature, providing not just the core model architectures but also the necessary tools and components for training, optimizing, and extending these models for diverse applications.

#### **4.7. Auto Classes (Prefix: Auto)**

A significant feature of the Hugging Face Transformers library is the set of 'Auto' classes, which greatly simplify the process of working with pre-trained models. These classes, including `AutoConfig`, `AutoModel`, `AutoTokenizer`, `AutoFeatureExtractor`, and `AutoImageProcessor`, act as dynamic loaders. Based on a provided pre-trained model identifier (which can be a name on the Hugging Face Hub or a local directory path), these classes automatically determine and instantiate the appropriate configuration, model, tokenizer, or processor.<sup>1</sup> The `from_pretrained()` method is central to these classes, handling the loading of the relevant components, including pre-trained



weights.<sup>1</sup> This abstraction eliminates the need for users to remember the specific class names associated with each model architecture, making it easier to experiment with different pre-trained models. For instance, to load a BERT model for sequence classification, a user can simply use

`AutoModelForSequenceClassification.from_pretrained('bert-base-uncased')` without needing to explicitly import `BertConfig`, `BertModel`, and `BertTokenizer`. The 'Auto' classes also facilitate the integration of custom models and configurations by providing registration mechanisms.<sup>1</sup> This allows users to extend the library with their own architectures while still benefiting from the convenience of the 'Auto' API. The following table summarizes the key 'Auto' classes and their functionalities:

Auto Class	Functionality
<code>AutoConfig</code>	Automatically loads the configuration for a given pre-trained model.
<code>AutoModel</code>	Automatically loads the base model for a given pre-trained model.
<code>AutoModelFor...</code>	Automatically loads a model with a specific task head (e.g., classification).
<code>AutoTokenizer</code>	Automatically loads the tokenizer associated with a given pre-trained model.
<code>AutoFeatureExtractor</code>	Automatically loads the feature extractor for a given pre-trained model (audio).
<code>AutoImageProcessor</code>	Automatically loads the image processor for a given pre-trained model (vision).

The 'Auto' classes represent a key design principle of the Hugging Face Transformers library, prioritizing ease of use and flexibility when working with a vast ecosystem of pre-trained models.

#### 4.8. Pipeline Classes (Suffix: Pipeline)

Pipeline classes, such as `AudioClassificationPipeline` and `AutomaticSpeechRecognitionPipeline`, offer a high-level, end-to-end abstraction for performing common machine learning tasks with minimal code.<sup>2</sup> These classes

encapsulate the entire workflow, from loading the appropriate pre-trained model and tokenizer/processor to preprocessing the input data, performing inference, and post-processing the model's output. Using a pipeline, a user can accomplish a complex task like sentiment analysis or audio classification in just a few lines of code, without needing to explicitly handle the individual steps of model loading, tokenization, and inference. For example, the `AudioClassificationPipeline` can take an audio file as input and directly output a classification label indicating the type of sound it contains. These pipeline classes make it incredibly easy for users, especially those new to the library or those seeking rapid prototyping, to leverage the power of pre-trained models for various applications. They abstract away the lower-level details, allowing users to focus on the task at hand rather than the intricacies of model loading and inference.

## 5. Illustrative Examples and Use Cases

To further illustrate the usage of the Hugging Face Transformers library and the roles of some of the listed classes, consider the following examples:

### Loading a Pre-trained Model and Tokenizer for Text Classification:

Python

```
from transformers import AutoModelForSequenceClassification, AutoTokenizer

# Load a pre-trained DistilBERT model fine-tuned for sentiment analysis
model_name = "distilbert-base-uncased-finetuned-sst-2-english"
model = AutoModelForSequenceClassification.from_pretrained(model_name)
tokenizer = AutoTokenizer.from_pretrained(model_name)

# Example usage
text = "This movie was absolutely amazing!"
inputs = tokenizer(text, return_tensors="pt")
outputs = model(**inputs)
# Process the outputs to get the sentiment label
```

This example demonstrates how `AutoModelForSequenceClassification` and `AutoTokenizer` can be used to easily load a pre-trained model and its corresponding tokenizer for the task of text classification.<sup>8</sup> The `from_pretrained()` method



automatically fetches the necessary configuration and weights from the Hugging Face Hub.

### Performing Inference with a Pipeline:

Python

```
from transformers import pipeline
```

```
# Create a sentiment analysis pipeline using the default model
sentiment_pipeline = pipeline("sentiment-analysis")
result = sentiment_pipeline("I am so happy today!")
print(result)
```

```
# Create a text generation pipeline with a specific model
text_generator = pipeline("text-generation", model="distilgpt2")
output = text_generator("The quick brown fox")
print(output)
```

This example showcases the simplicity of using the pipeline function for performing inference on common tasks like sentiment analysis and text generation.<sup>8</sup> The pipeline function automatically selects an appropriate pre-trained model for the specified task if one is not explicitly provided.

### Tokenizing Text Data and Obtaining Input IDs:

Python

```
from transformers import AutoTokenizer
```

```
# Load the tokenizer for a pre-trained BERT model
tokenizer = AutoTokenizer.from_pretrained("bert-base-uncased")
```

```
# Tokenize a sentence
sentence = "Hello, world!"
```

```
inputs = tokenizer(sentence)
print(inputs)
print("Input IDs:", inputs["input_ids"])
print("Attention Mask:", inputs["attention_mask"])
```

This example illustrates how AutoTokenizer can be used to tokenize a sentence and obtain the numerical input\_ids and attention\_mask that are required as input for a transformer model.<sup>22</sup>

### Loading Configuration and Model Separately:

Python

```
from transformers import AutoConfig, AlbertModel

# Load the configuration for an Albert model
config = AutoConfig.from_pretrained("albert-base-v2")

# Initialize an Albert model from the configuration (without pre-trained weights)
model = AlbertModel(config)

print(model.config)
```

This example demonstrates how to load a model's configuration using AutoConfig and then use that configuration to initialize a model architecture using the base model class (AlbertModel).<sup>27</sup> This approach can be useful when one wants to initialize a model with a specific architecture before potentially loading pre-trained weights or training from scratch.

## 6. Conclusion: Understanding the Architecture and Functionality of Hugging Face Transformers

In conclusion, the Hugging Face Transformers library stands as a powerful and versatile tool for developers and researchers seeking to leverage the capabilities of pre-trained machine learning models. Its extensive collection of models, support for multiple deep learning frameworks, and user-friendly API have made it a cornerstone in the field. Understanding the different categories of classes within the library – Configuration, Model, Tokenizer, Processor, Auto, Pipeline, and Utility – is crucial for

effectively navigating its vast functionalities. Configuration classes define model architectures, while model classes implement them, often with task-specific extensions. Tokenizer classes handle the essential conversion of text to numerical inputs, with 'Fast' variants offering performance optimizations. Processor classes are vital for preparing data for multimodal models, and feature extractors/image processors handle non-textual data. The 'Auto' classes significantly simplify the process of loading pre-trained models and associated components, and pipeline classes provide high-level abstractions for common machine learning tasks. The library's integration with the Hugging Face Model Hub further enhances its value by providing a central platform for discovering and sharing models and datasets. By providing access to state-of-the-art pre-trained models and the necessary tools to utilize them effectively, the Hugging Face Transformers library democratizes access to advanced machine learning, fostering innovation and accelerating progress across various domains. Continued exploration of the official documentation is highly recommended for users to fully leverage the specific models and functionalities relevant to their individual needs.

## Works cited

1. Auto Classes - Hugging Face, accessed on April 29, 2025, [https://huggingface.co/docs/transformers/en/model\\_doc/auto](https://huggingface.co/docs/transformers/en/model_doc/auto)
2. AutoModels — transformers 3.0.2 documentation - Hugging Face, accessed on April 29, 2025, [https://huggingface.co/transformers/v3.0.2/model\\_doc/auto.html](https://huggingface.co/transformers/v3.0.2/model_doc/auto.html)
3. Transformers - Hugging Face, accessed on April 29, 2025, <https://huggingface.co/docs/transformers/v4.13.0/index>
4. Transformers - Hugging Face, accessed on April 29, 2025, <https://huggingface.co/docs/transformers/v4.36.1/index>
5. Transformers - Hugging Face, accessed on April 29, 2025, <https://huggingface.co/docs/transformers/index>
6. Transformers: State-of-the-art Machine Learning for Pytorch, TensorFlow, and JAX. - GitHub, accessed on April 29, 2025, <https://github.com/huggingface/transformers>
7. transformers 3.0.2 documentation - Hugging Face, accessed on April 29, 2025, <https://huggingface.co/transformers/v3.0.2/index.html>
8. Using transformers at Hugging Face, accessed on April 29, 2025, <https://huggingface.co/docs/hub/transformers>
9. Documentation - Hugging Face, accessed on April 29, 2025, <https://huggingface.co/docs>
10. Change package name from "transformers" to something less generic #24934 - GitHub, accessed on April 29, 2025, <https://github.com/huggingface/transformers/issues/24934>
11. Quickstart - Hugging Face, accessed on April 29, 2025, <https://huggingface.co/docs/transformers/quicktour>

12. Auto Classes - Hugging Face, accessed on April 29, 2025,  
[https://huggingface.co/docs/transformers/v4.21.2/model\\_doc/auto](https://huggingface.co/docs/transformers/v4.21.2/model_doc/auto)
13. Configuration - Hugging Face, accessed on April 29, 2025,  
[https://huggingface.co/docs/transformers/main\\_classes/configuration](https://huggingface.co/docs/transformers/main_classes/configuration)
14. transformers/src/transformers/models/auto/configuration\_auto.py at main - GitHub, accessed on April 29, 2025,  
[https://github.com/huggingface/transformers/blob/main/src/transformers/models/auto/configuration\\_auto.py](https://github.com/huggingface/transformers/blob/main/src/transformers/models/auto/configuration_auto.py)
15. Customizing models - Hugging Face, accessed on April 29, 2025,  
[https://huggingface.co/docs/transformers//custom\\_models](https://huggingface.co/docs/transformers//custom_models)
16. Models - Hugging Face, accessed on April 29, 2025,  
[https://huggingface.co/docs/transformers/v4.51.3/en/main\\_classes/model](https://huggingface.co/docs/transformers/v4.51.3/en/main_classes/model)
17. What is the classification head doing exactly? - Transformers - Hugging Face Forums, accessed on April 29, 2025,  
<https://discuss.huggingface.co/t/what-is-the-classification-head-doing-exactly/10138>
18. What are differences between AutoModelForSequenceClassification vs AutoModel, accessed on April 29, 2025,  
<https://stackoverflow.com/questions/69907682/what-are-differences-between-a-automodelforsequenceclassification-vs-automodel>
19. Difference BertModel, AutoModel and AutoModelForMaskedLM - Hugging Face Forums, accessed on April 29, 2025,  
<https://discuss.huggingface.co/t/difference-bertmodel-automodel-and-automodelformaskedlm/24340>
20. How Labelled Data is Processed | Transformers Trainer - Hugging Face Forums, accessed on April 29, 2025,  
<https://discuss.huggingface.co/t/how-labelled-data-is-processed-transformers-trainer/80644>
21. Models - Hugging Face, accessed on April 29, 2025,  
[https://huggingface.co/docs/transformers/main/main\\_classes/model](https://huggingface.co/docs/transformers/main/main_classes/model)
22. Tokenizers - Hugging Face, accessed on April 29, 2025,  
[https://huggingface.co/docs/transformers/main/fast\\_tokenizers](https://huggingface.co/docs/transformers/main/fast_tokenizers)
23. Tokenizer - Hugging Face, accessed on April 29, 2025,  
[https://huggingface.co/docs/transformers/main\\_classes/tokenizer](https://huggingface.co/docs/transformers/main_classes/tokenizer)
24. Summary of the tokenizers - Hugging Face, accessed on April 29, 2025,  
[https://huggingface.co/docs/transformers/tokenizer\\_summary](https://huggingface.co/docs/transformers/tokenizer_summary)
25. Putting it all together - Hugging Face LLM Course, accessed on April 29, 2025,  
<https://huggingface.co/learn/llm-course/chapter2/6>
26. How to pass possible class names to distilbert - Stack Overflow, accessed on April 29, 2025,  
<https://stackoverflow.com/questions/68800382/how-to-pass-possible-class-names-to-distilbert>
27. Two questions when I wrapped the AutoModelForMaskedLM - Transformers, accessed on April 29, 2025,  
<https://discuss.huggingface.co/t/two-questions-when-i-wrapped-the-automodelf>

[ormaskedlm/146555](#)

28. How to use gated models? - Hub - Hugging Face Forums, accessed on April 29, 2025, <https://discuss.huggingface.co/t/how-to-use-gated-models/53234>